

Prometheus gRPC errors graphics

Status

View alerts

✓Cluster

✓Control Plane

✓Operators

✓Insights

0 issues found

✓Dynamic Plugins

vSphere connection

No errors are reported. Click the link for details.

!etcdGRPCRequestsSlow

View details

20 de nov. de 2024, 16:01

etcd cluster "etcd": 99th percentile of gRPC requests is 10s on etcd instance 172.23.218.103:9979 for MemberList method.

!KubeAPIErrorBudgetBurn

View details

20 de nov. de 2024, 15:43

The API server is burning too much error budget. This alert fires when too many requests are failing with high latency. Use the 'API Performance' monitoring dashboards to narrow down the request states and latency. The 'etcd' monitoring dashboards also provides metrics to help determine etcd stability and performance.

!etcdGRPCRequestsSlow

View details

20 de nov. de 2024, 15:41

etcd cluster "etcd": 99th percentile of gRPC requests is 9.728729281767956s on etcd

Status

View alerts

Cluster

✓Control Plane

✓Operators

✓Insights

0 issues found

✓Dynamic Plugins

!vSphere connection

Missing the vSphere config map.

!etcdGRPCRequestsSlow

View details

20 de nov. de 2024, 16:13

etcd cluster "etcd": 99th percentile of gRPC requests is 10s on etcd instance 172.23.218.104:9979 for MemberList method.

!etcdGRPCRequestsSlow

View details

20 de nov. de 2024, 16:13

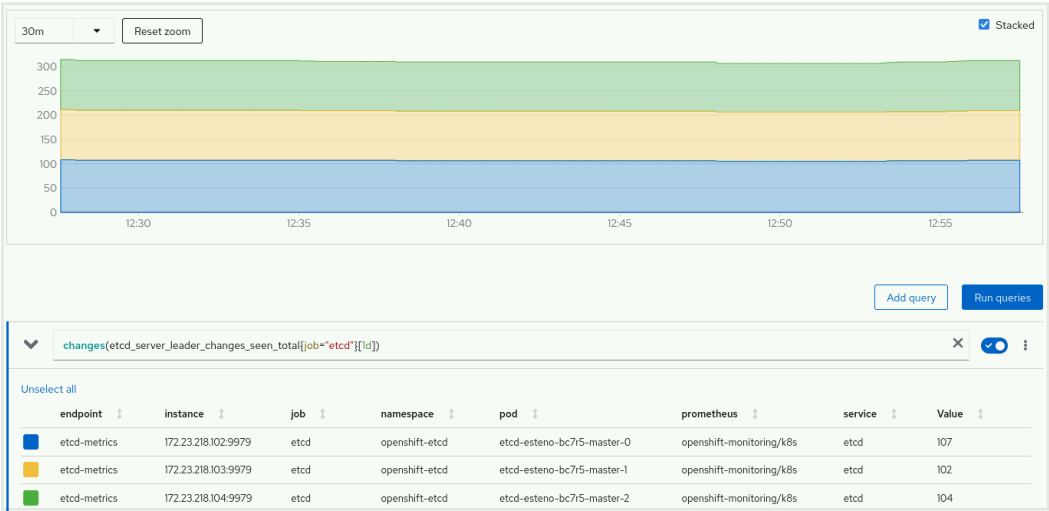
etcd cluster "etcd": 99th percentile of gRPC requests is 9.03354430379747s on etcd instance 172.23.218.102:9979 for Txn method.

!etcdGRPCRequestsSlow

View details

20 de nov. de 2024, 16:13

etcd cluster "etcd": 99th percentile of gRPC requests is 10s on etcd instance 172.23.218.102:9979 for Range method.



Name
KubeAPIErrorBudgetBurn

Severity
Critical

Description
The API server is burning too much error budget. This alert fires when too many requests are failing with high latency. Use the 'API Performance' monitoring dashboards to narrow down the request states and latency. The 'etcd' monitoring dashboards also provides metrics to help determine etcd stability and performance.

Summary
The API server is burning too much error budget.

Runbook
<https://github.com/openshift/runbooks/blob/master/alerts/cluster-kube-apiserver-operator/KubeAPIErrorBudgetBurn.md>

Labels
long=6h namespace=openshift-kube-apiserver prometheus=openshift-monitoring/k8s severity=critical short=30m alertname=KubeAPIErrorBudgetBurn

Alerting rule
KubeAPIErrorBudgetBurn

Source
Platform

State
Firing
Since 21 de nov. de 2024, 12:10